

HR Analytics Project - Understanding the Attrition in HR

Problem Statement:-

Every year a lot of companies hire a number of employees. The companies invest time and money in training those employees, not just this but there are training programs within the companies for their existing employees as well.

The objective of the model to increase the effectiveness of their employees and reduce the time and money investing in employees.

HR Analytics:-

Human resource analytics (HR analytics) is an area in the field of analytics that refers to applying analytic processes to the human resource department of an organization in the hope of improving employee performance and therefore getting a better return on investment. HR analytics does not just deal with gathering data on employee efficiency. Instead, it aims to provide insight into each process by gathering data and then using it to make relevant decisions about how to improve these processes.

Attrition in HR:-

Attrition in human resources refers to the gradual loss of employee's overtime. In general, relatively high attrition is problematic for companies. HR professionals often assume a leadership role in designing company compensation programs, work culture, and motivation systems that help the organization retain top employees.

How does Attrition affect companies? And how does HR Analytics help in analyzing attrition? We will discuss the first question here and for the second question, we will write the code and try to understand the process step by step.

Attrition affecting Companies:-

A major problem in high employee attrition is its cost to an organization. Job postings, hiring processes, paperwork, and new hire training are some of the common expenses of losing employees and replacing them. Additionally, regular employee turnover prohibits your organization from increasing its collective knowledge base and experience over time. This is especially concerning if your business is customer-facing, as customers often prefer to interact with familiar people. Errors and issues are more likely if you constantly have new workers.

Importing the Libraries:-

```
import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import warnings
warnings.filterwarnings('ignore')
from matplotlib import pyplot as plt
```

EXPLORATORY DATA ANALYSIS:-

```
In [4]: df.dtypes

Out[4]: Age                int64
Attrition                object
BusinessTravel            object
DailyRate                int64
Department                object
DistanceFromHome          int64
Education                 int64
EducationField            object
EmployeeCount             int64
EmployeeNumber            int64
EnvironmentSatisfaction   int64
Gender                    object
HourlyRate                int64
JobInvolvement            int64
JobLevel                  int64
JobRole                   object
JobSatisfaction           int64
MaritalStatus             object
MonthlyIncome             int64
MonthlyRate               int64
NumCompaniesWorked        int64
Over18                    object
OverTime                  object
PercentSalaryHike         int64
PerformanceRating         int64
RelationshipSatisfaction   int64
StandardHours             int64
StockOptionLevel          int64
TotalWorkingYears         int64
TrainingTimesLastYear     int64
WorkLifeBalance           int64
YearsAtCompany            int64
YearsInCurrentRole        int64
YearsSinceLastPromotion    int64
YearsWithCurrManager       int64
dtype: object
```

The features 'Attrition','BusinessTravel','Department','EducationField','Gender','JobRole','MaritalStatus','OverTime',over18 are object type remaining all are integer type

Numeric variables:-

- Related to personal information: age, distance_from_home, employee_number
- Related to income: hourly_rate, daily_rate, monthly_rate, monthly_income, percent_salary_hike

Related to duration in company: years_at_company, years_in_current_role, years_since_last_promotion, years_with_curr_manager, total_working_years, num_companies_worked, standard_hours, training_times_last_year, employee_count

Categorical variables:-

- Binary variables: attrition(target variable), gender, over18, over_time
- Nominal variables: department, education_field, job_role, marital_status
- Ordinal variables:

Ordinal regarding satisfaction and performance: environment_satisfaction, job_satisfaction, relationship_satisfaction, work_life_balance, job_involvement, performance_rating

➤ Other ordinal: business travel, education, job_level, stock_option_level

```
In [5]: df.isnull().sum()

Out[5]: Age                0
Attrition                 0
BusinessTravel            0
DailyRate                0
Department               0
DistanceFromHome         0
Education                0
EducationField            0
EmployeeCount            0
EmployeeNumber           0
EnvironmentSatisfaction  0
Gender                   0
HourlyRate               0
JobInvolvement           0
JobLevel                 0
JobRole                  0
JobSatisfaction          0
MaritalStatus            0
MonthlyIncome            0
MonthlyRate              0
NumCompaniesWorked       0
Over18                   0
OverTime                 0
PercentSalaryHike        0
PerformanceRating        0
RelationshipSatisfaction  0
StandardHours            0
StockOptionLevel         0
TotalWorkingYears        0
TrainingTimesLastYear    0
WorkLifeBalance          0
YearsAtCompany           0
YearsInCurrentRole       0
YearsSinceLastPromotion  0
YearsWithCurrManager     0
dtype: int64

dataset has no null values
```

Activate Window
Go to Settings to activate window

This dataset has no null values

```
In [6]: df['StandardHours'].unique


Out[6]: <bound method Series.unique of 0      80
1      80
2      80
3      80
4      80
..
1465   80
1466   80
1467   80
1468   80
1469   80
Name: StandardHours, Length: 1470, dtype: int64>

In [7]: df['EmployeeCount'].unique

Out[7]: <bound method Series.unique of 0      1
1      1
2      1
3      1
4      1
..
1465   1
1466   1
1467   1
1468   1
1469   1
Name: EmployeeCount, Length: 1470, dtype: int64>
```

Displaying value count of unique value in each feature. To identify column having single unique value

UNI VARIATE ANALYSIS:-



The features standardhours,over18 and employeecount has only one value so it wont create any impact on the target feature Attrition.


The features standard hours, over18 and employee count has only one value so it won't create any impact on the target feature Attrition.

All the three columns having single value so I'm going to dropping it from the given dataset.

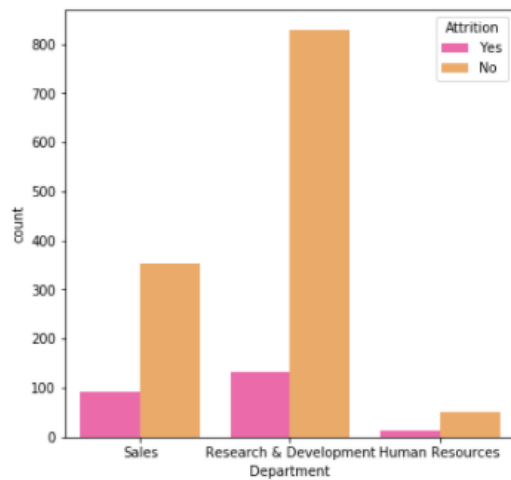
BI VARIATE ANALYSIS (Categorical columns vs Target)

```
In [12]: plt.figure(figsize=(10,5))
sns.countplot(df.BusinessTravel,palette="Set1",hue=df.Attrition)

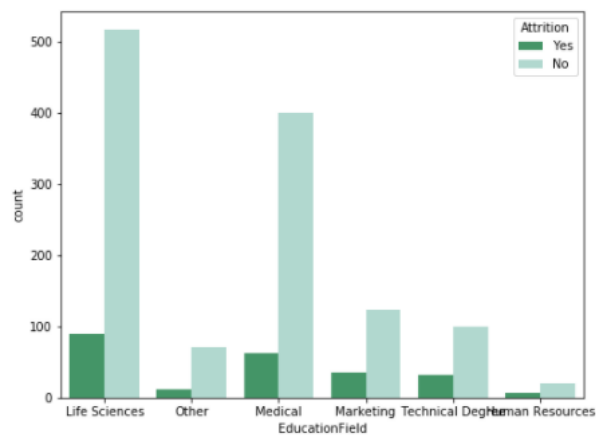
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x20b09313708>
```




From the above chart it is concluded that employees who travel rarely have high attrition.




From the above plot it's apparent that comparatively the employees who belongs to Research and Development will like to continue their job.



From the above plot it's apparent that comparatively the employees who belongs to both lifesciece and medical field are not resigning their job.



Comparing the percentage of attrition out of 588 female only 88 people are quitting.




Distance from home is not an important feature to create impact on Attrition feature. Distance from home is not impact on gender.

EDA CONCLUSION:-

Employees who belongs to below category having less attrition on rate

- travel rarely
- who belongs to R&D department
- who belongs to life science and medical field
- female
- working as sales executive and research scientists
- unmarried
- not working over time
- moderate work life balance
- high job involvement
- working in single company
- performance rating:3
- employees who having 0 stocks
- low monthly income.

We cannot predict using relationship satisfaction, job satisfaction, Environment satisfaction features




Over time feature is highly correlated with attrition.

Handling outliers in numerical column:-


```
In [28]: sns.boxplot(df['MonthlyIncome'],color="cyan")
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x20b099b27c8>
```



```
In [29]: sns.boxplot(df['MonthlyRate'],color="blue")
```

```
Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x20b0993ddc8>
```



```
In [32]: z1 = np.abs(stats.zscore(df_new['MonthlyIncome']))
print(z1)
```

```
[0.10834951 0.29171859 0.93765369 ... 0.07669019 0.23647414 0.44597809]
```

```
In [33]: df_new['MonthlyIncome'] = df_new.MonthlyIncome[(z1<3)]
df_new.shape
```

```
Out[33]: (1470, 32)
```

outliers are removed from numerical data monthly income

Data pre-processing:-

The features standard hours, over18 and employee count has only single value so it won't create any impact on the target feature Attrition.

Employee Number feature is just an identifier and it's not required for modelling either. So I'm dropping these features

The features standardhours,over18 and employeecount has only one value so it wont create any impact on the target feature Attrition.

```
In [4]: 1 cols=['StandardHours','Over18','EmployeeCount']
        2 df_new=df.drop(cols,axis=1)
```

All the three columns having single value so I'm dropping it from the given dataset

```
In [4]: 1 df_new.shape
```

```
Out[4]: (1470, 32)
```


Encoding all categorical column into numerical column using label encoding technique

Encoding all categorical column into numerical column using label encoding technique

```
In [36]: data_clean=df_new
col_encod=['Attrition','BusinessTravel','Department','EducationField','Gender','JobRole','MaritalStatus','OverTime']

In [37]: from sklearn import preprocessing
for col in col_encod:
    label = preprocessing.LabelEncoder()
    data_clean[col]= label.fit_transform(df_new[col])
```

CHECKING FOR SKEWNESS AND USING METHOD TO REMOVE SKEWNESS:-

```
In [40]: data_clean.skew()
```

```
Out[40]: Age                0.413286
Attrition                1.844366
BusinessTravel          -1.439006
DailyRate              -0.003519
Department              0.172231
DistanceFromHome        0.958118
Education              -0.289681
EducationField          0.550371
EmployeeNumber          0.016574
EnvironmentSatisfaction -0.321654
Gender                 -0.408665
HourlyRate             -0.032311
JobInvolvement          -0.498419
JobLevel                1.025401
JobRole                -0.357270
JobSatisfaction         -0.329672
MaritalStatus           -0.152175
MonthlyIncome           1.369817
MonthlyRate             0.018578
NumCompaniesWorked      1.026471
OverTime                0.964489
PercentSalaryHike       0.821128
PerformanceRating       1.921883
RelationshipSatisfaction -0.302828
StockOptionLevel        0.968980
TotalWorkingYears       1.117172
TrainingTimesLastYear   0.553124
WorkLifeBalance         -0.552480
YearsAtCompany          1.764529
YearsInCurrentRole      0.917363
YearsSinceLastPromotion 1.984290
YearsWithCurrManager    0.833451
dtype: float64
```

```
In [41]: from sklearn.preprocessing import power_transform
x = power_transform(data_clean,method='yeo-johnson')
```

We can check that our input dataset X is having some skewness so we used yeo-johnson method to remove the skewness from our dataset

HANDLING CLASS IMBALANCE:-

Classification problem where the distribution of examples across the known classes is biased or skewed. To avoid this we are using SMOTE technique

SMOTE synthetic over-sampling works to cause the classifier to build larger decision regions that contain nearby minority class points. This will in turn avoid data loss.

```
In [42]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x = sc.fit_transform(x)
```

As we have checked earlier our dataset having class imbalancing so we are using standard scaler to balance the dataset before initializing building of our model.

MODELLING:-

It is a binary classification problem so I have modelled using logistic regression and other classification models.

```
In [45]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import cross_val_score
from matplotlib import pyplot
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
x=scaled
y=y1
```

```
In [46]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.25,random_state=1)

models=[LogisticRegression(),RandomForestClassifier(),KNeighborsClassifier(),GradientBoostingClassifier(),DecisionTreeClassifier()]
scorelist=[]
acclist=[]
```

```
In [47]: def create_model(model):
m=model
m.fit(xtrain,ytrain)
p=m.predict(xtest)
score=m.score(xtest,ytest)
result = confusion_matrix(ytest,p)
result1 = classification_report(ytest,p)
result2 = accuracy_score(ytest,p)
scorelist.append(score)
acclist.append(result2)

print(m,"\n")

print('Accuracy score:',score,"\n")
print("Confusion Matrix:\n",result)
print('classification_report\n',result1)
print('Average accuracy_score',result2)
print('-----')
for i in models:
create_model(i)

print('Maximun accuracy Score is shown by',models[acclist.index(max(acclist))],max(acclist))
```

LogisticRegression()

Accuracy score: 0.8586956521739131

"Confusion Matrix:

" [[297 3]
[49 19]]

classification_report

	precision	recall	f1-score	support
0	0.86	0.99	0.92	300
1	0.86	0.28	0.42	68

accuracy			0.86	368
macro avg	0.86	0.63	0.67	368
weighted avg	0.86	0.86	0.83	368

Average accuracy_score 0.8586956521739131

RandomForestClassifier()

Accuracy score: 0.8342391304347826

"Confusion Matrix:

" [[296 4]
[57 11]]

classification_report

	precision	recall	f1-score	support
0	0.84	0.99	0.91	300
1	0.73	0.16	0.27	68

accuracy			0.83	368
macro avg	0.79	0.57	0.59	368
weighted avg	0.82	0.83	0.79	368

Average accuracy_score 0.8342391304347826

KNeighborsClassifier()

Accuracy score: 0.8315217391304348

"Confusion Matrix:

" [[295 5]
[57 11]]

classification_report

	precision	recall	f1-score	support
0	0.84	0.98	0.90	300
1	0.69	0.16	0.26	68

accuracy			0.83	368
macro avg	0.76	0.57	0.58	368
weighted avg	0.81	0.83	0.79	368

Average accuracy_score 0.8315217391304348

GradientBoostingClassifier()

Accuracy score: 0.8288043478260869

"Confusion Matrix:

" [[286 14]
[49 19]]

classification_report

	precision	recall	f1-score	support
0	0.85	0.95	0.90	300
1	0.58	0.28	0.38	68

accuracy			0.83	368
macro avg	0.71	0.62	0.64	368
weighted avg	0.80	0.83	0.80	368

Average accuracy_score 0.8288043478260869

Maximum accuracy Score is shown by LogisticRegression () 0.8586956521739131

Cross Validation:-

In order to avoid over fitting, Cross-validation is used to estimate the skill of a machine learning model on unseen data.

```
In [49]: score1=[]
```

```
In [50]: lr=LogisticRegression()  
scores=cross_val_score(lr,x,y,cv=5)  
score1.append(scores)  
scores
```

```
Out [50]: array([0.87755102, 0.86734694, 0.86054422, 0.8707483 , 0.87755102])
```

```
In [51]: rf=RandomForestClassifier()  
scores=cross_val_score(rf,x,y,cv=5)  
score1.append(scores)  
scores
```

```
Out [51]: array([0.8537415 , 0.85714286, 0.86054422, 0.86054422, 0.85714286])
```

```
In [52]: kn=KNeighborsClassifier()  
scores=cross_val_score(kn,x,y,cv=5)  
score1.append(scores)  
scores
```

```
Out [52]: array([0.83673469, 0.83673469, 0.8537415 , 0.85034014, 0.84693878])
```

```
In [53]: gbt=GradientBoostingClassifier()  
scores=cross_val_score(gbt,x,y,cv=5)  
score1.append(scores)  
scores
```

```
Out [53]: array([0.84693878, 0.87414966, 0.87414966, 0.85714286, 0.86054422])
```

```
In [54]: dt=DecisionTreeClassifier()  
scores=cross_val_score(dt,x,y,cv=5)  
score1.append(scores)  
scores
```

```
Out [54]: array([0.7755102 , 0.78911565, 0.82312925, 0.76870748, 0.77891156])
```

Difference of predicted model and crossvalidation score

```
In [55]: models=[LogisticRegression(),RandomForestClassifier(),KNeighborsClassifier(),GradientBoostingClassifier(),DecisionTreeClassifier()]
for i in range(0,5):
    print(models[i],"difference is",scorel[i]-acclist[i])
```

```
LogisticRegression() difference is [0.01885537 0.00865129 0.00184857 0.01205265 0.01885537]
RandomForestClassifier() difference is [0.01950237 0.02290373 0.02630509 0.02630509 0.02290373]
KNeighborsClassifier() difference is [0.00521295 0.00521295 0.02221976 0.0188184 0.01541704]
GradientBoostingClassifier() difference is [0.01813443 0.04534531 0.04534531 0.02833851 0.03173987]
DecisionTreeClassifier() difference is [0.02279281 0.03639825 0.07041186 0.01599009 0.02619417]
```

from the observation KNeighborsClassifier() has least difference so I'm selecting KNeighborsClassifier() as best model

From the observation KNeighborsClassifier () has least difference so I'm selecting KNeighborsClassifier () as best model.

Hyper Tuning:

```
In [56]: from sklearn.model_selection import GridSearchCV,KFold
params = {
    'n_neighbors' : [5,7,9,11,13,15],
    'weights' : ['uniform','distance'],
    'metric' : ['minkowski','euclidean','manhattan'],
    'p':[1,2], 'leaf_size':list(range(1,20))
}

gs2 = GridSearchCV(KNeighborsClassifier(), params, verbose = 1, cv=3, n_jobs = -1)
gs2.fit(xtrain, ytrain)
print('Best param:', gs2.best_params_)

Fitting 3 folds for each of 1368 candidates, totalling 4104 fits
Best param: {'leaf_size': 1, 'metric': 'minkowski', 'n_neighbors': 5, 'p': 2, 'weights': 'uniform'}
```

Best parameters: {'leaf_size': 1, 'metric': 'minkowski', 'n_neighbors': 5, 'p': 1, 'weights': 'distance'}

Modelling using best parameter and best model:-

```
In [57]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=.25)
model = KNeighborsClassifier(metric='minkowski', n_neighbors=5, weights='distance', p=1, leaf_size=1)
model.fit(x_train, y_train)
model.score(x_test, y_test)

Out[57]: 0.8505434782608695

In [58]: y_pred_1 = model.predict(x_test)

In [59]: result = confusion_matrix(y_test, y_pred_1)
print("Confusion Matrix:")
print(result)
result1 = classification_report(y_test, y_pred_1)
print("Classification Report:")
print(result1)
result2 = accuracy_score(y_test, y_pred_1)
print("Accuracy:", result2)

Confusion Matrix:
[[306  1]
 [ 54  71]]
Classification Report:
              precision    recall  f1-score   support

     0           0.85       1.00       0.92         307
     1           0.88       0.11       0.20          61

   accuracy          0.86
  macro avg          0.86
 weighted avg          0.85

Accuracy: 0.8505434782608695
```


Final model after hyper tuning with accuracy 0.8586956521739131

Best model: KNeighbourClassifier Best param: {'leaf_size': 1, 'metric': 'minkowski', 'n_neighbors': 5, 'p': 1, 'weights': 'distance'}

ROC AUC CURVE:-

```
In [60]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.26,random_state=1)
m=KNeighborsClassifier(metric='minkowski', n_neighbors=5,weights='distance',p=1,leaf_size=1)
m.fit(xtrain,ytrain)
p=m.predict(xtest)
ns_probs = [0 for _ in range(len(ytest))]
m_probs = p
ns_auc = roc_auc_score(ytest, ns_probs)
m_auc = roc_auc_score(ytest, m_probs)
print('No Skill: ROC AUC=%3f' % (ns_auc))
print('model: ROC AUC=%3f' % (m_auc))
ns_fpr, ns_tpr,_= roc_curve(ytest, ns_probs)
m_fpr, m_tpr,_=roc_curve(ytest, m_probs)
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(m_fpr, m_tpr, marker='.', label='Randomforest')
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
pyplot.legend()
pyplot.show()
```

No Skill: ROC AUC=0.500
model: ROC AUC=0.557



Conclusion:

I have developed a model to predict attrition of an employee with 85.8% accuracy

Saving the model:-

```
In [61]: from joblib import dump
dump(model, 'model_hr.joblib')
```

```
Out[61]: ['model_hr.joblib']
```

```
In [62]: from joblib import load
loaded = load('model_hr.joblib')
```

```
In [ ]:
```